

Linux Foundation

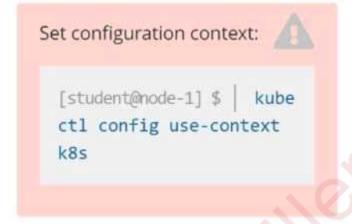
CKA Exam

Linux Foundation Certified Kubernetes Administrator Exam

SIMULATION

Monitor the logs of pod foo and:

- Extract log lines corresponding to error unable-to-access-website
- Write them to /opt/KULM00201/foo



Solution



SIMULATION

List all persistent volumes sorted by capacity, saving the full kubectl output to /opt/KUCC00102/volume_list. Use kubectl 's own functionality for sorting the output, and do not manipulate it any further.

Solution

solution

Readr	ne	_ Web Terminal		LITHELI	NUXFOUNDATION
77d pv0007	7Gi	RWO	Recycle	Available	slow
77d	761	KWO	recycle	Available	SIOW
pv0006 77d	8Gi	RWO	Recycle	Available	slow
pv0003 77d	10Gi	RWO	Recycle	Available	slow
77d	11Gi	RWO	Recycle	Available	slow
77d	13Gi	RWO	Recycle	Available	slow
77d	14Gi	RWO	Recycle	Available	slow
77d	16Gi	RWO	Recycle	Available	slow
77d	17Gi	RWO	Recycle	Available	slow
77d	18Gi	RWO	Recycle	Available	slow
77d	19Gi	RWO	Recycle	Available	slow
77d	21Gi	RWO	Recycle	Available	slow JCC00102/volume list

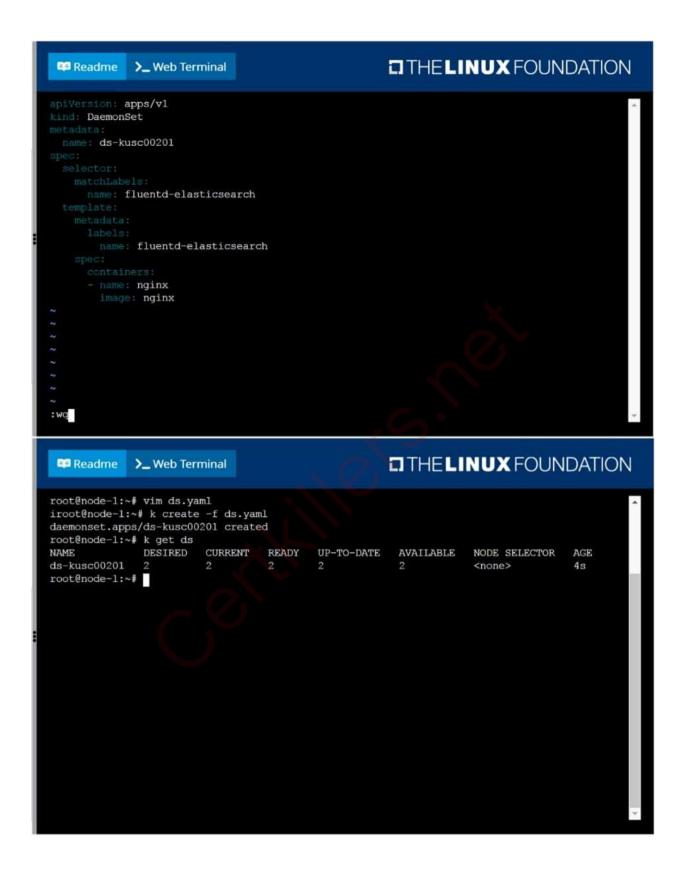
Question: 3

SIMULATION

Ensure a single instance of pod nginx is running on each node of the Kubernetes cluster where nginx also represents the Image name which has to be used. Do not override any taints currently in place. Use DaemonSet to complete this task and use ds-kusc00201 as DaemonSet name.

Solution





SIMULATION

Perform the following tasks:

• Add an init container to hungry-bear (which has been defined in spec file /opt/KUCC00108/pod-spec-KUC

C00108.yaml

)

- The init container should create an empty file named /workdir/calm.txt
- If /workdir/calm.txt is not detected, the pod should exit
- Once the spec file has been updated with the init container definition, the pod should be created

Solution



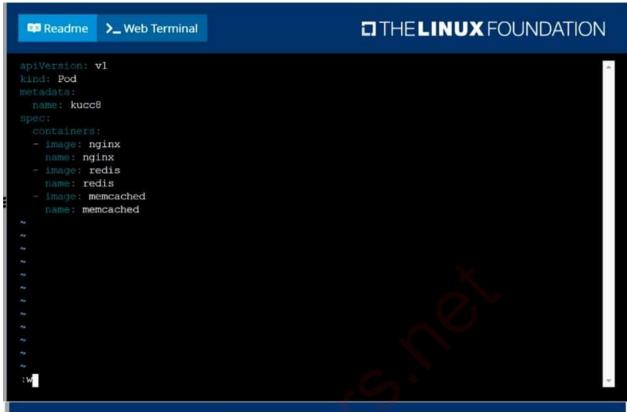


SIMULATION

Create a pod named kucc8 with a single app container for each of the following images running inside (there may be between 1 and 4 images specified): nginx + redis + memcached.

Solution

```
THE LINUX FOUNDATION
 Readme
             >_ Web Terminal
root@node-1:~# vim ds.yaml
iroot@node-1:~# k create -f ds.yaml
daemonset.apps/ds-kusc00201 created
root@node-1:~# k get ds
                         CURRENT
                                           UP-TO-DATE
                                                        AVAILABLE
                                                                    NODE SELECTOR
               DESIRED
                                   READY
ds-kusc00201
                                                                    <none>
                                                                                    43
root@node-1:~# vim /opt/KUCC00108/pod-spec-KUCC00108.yaml
root@node-1:~# k create -f /opt/KUCC00108/pod-spec-KUCC00108.yaml
pod/hungry-bear created
root@node-1:~# k get po
NAME
                             READY
                                     STATUS
                                               RESTARTS
                                                          AGE
cpu-utilizer-98b9se
                                                          5h50m
                             1/1
                                     Running
                                               0
                             1/1
cpu-utilizer-ab2d3s
                                     Running
                                                          5h50m
cpu-utilizer-kipb9a
                             1/1
                                                          5h50m
                                     Running
                                               0
ds-kusc00201-2r2k9
                             1/1
                                     Running
                                                          4m50s
ds-kusc00201-hzm9q
                             1/1
                                     Running
                                                          4m50s
                                                          5h52m
                             1/1
foo
                                     Running
                                               0
                                                          5h52m
front-end
                             1/1
                                     Running
hungry-bear
                             1/1
                                     Running
                                               0
                                                          42s
webserver-84c55967f4-qzjcv
                             1/1
                                                          6h7m
                                     Running
webserver-84c55967f4-t4791
                             1/1
                                               0
                                                          6h7m
                                     Running
root@node-1:~# k run nginx --image=nginx --dry-run=client -o yaml > nginx.yaml
root@node-1:~# vim nginx.yaml
```



Readme >_ Web Termina	al			1 IHEL	INUXFOUND	ATION
cpu-utilizer-98b9se	1/1	Running		0	5h51m	
pu-utilizer-ab2d3s	1/1	Running		0	5h51m	
cpu-utilizer-kipb9a	1/1	Running		0	5h51m	
ds-kusc00201-2r2k9	1/1	Running		0	6m12s	
ds-kusc00201-hzm9q	1/1	Running		0	6m12s	
foo	1/1	Running		0	5h54m	
front-end	1/1	Running		0	5h53m	
hungry-bear	1/1	Running		0	2m4s	
kucc8	0/3	Container	rCreating	0	4s	
webserver-84c55967f4-qzjcv	1/1	Running		0	6h9m	
webserver-84c55967f4-t4791	1/1	Running		0	6h9m	
root@node-1:~# k get po						
NAME	READY	STATUS	RESTARTS	AGE		
cpu-utilizer-98b9se	1/1	Running	0	5h52m		
cpu-utilizer-ab2d3s	1/1	Running	0	5h52m		
cpu-utilizer-kipb9a	1/1	Running	0	5h52m		
ds-kusc00201-2r2k9	1/1	Running	0	6m31s		
ds-kusc00201-hzm9q	1/1	Running	0	6m31s		
foo	1/1	Running	0	5h54m		
front-end	1/1	Running	0	5h54m		
hungry-bear	1/1	Running	0	2m23s		
kucc8	3/3	Running	0	23s		
webserver-84c55967f4-qzjcv	1/1	Running	0	6h9m		
webserver-84c55967f4-t4791	1/1	Running	0	6h9m		
root@node-1:~#						

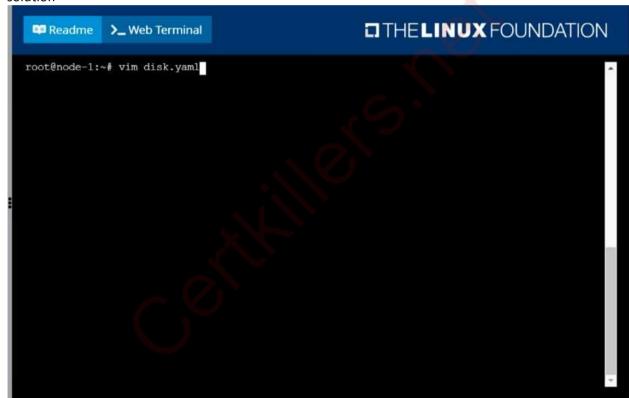
SIMULATION

Schedule a pod as follows:
• Name: nginx-kusc00101

• Image: nginx

• Node selector: disk=ssd

Solution





SIMULATION

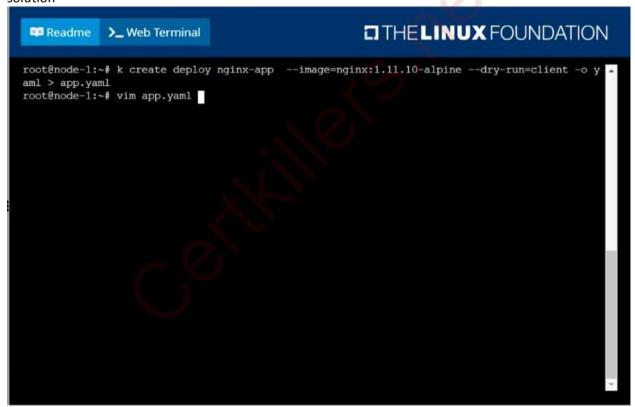
Create a deployment as follows:

- Name: nginx-app
- Using container nginx with version 1.11.10-alpine
- The deployment should contain 3 replicas

Next, deploy the application with new version 1.11.13-alpine, by performing a rolling update.

Finally, rollback that update to the previous version 1.11.10-alpine.

Solution



```
apiVersion: apps/v1
kind: Deployment
metadata:
name: nginx-app
spec:
replicas: 3
selector:
matchiabels:
app: nginx-app
template:
metadata:
labels:
app: nginx-app
spec:
containers:
- inage: nginx:1.11.10-alpine
name: nginx
```

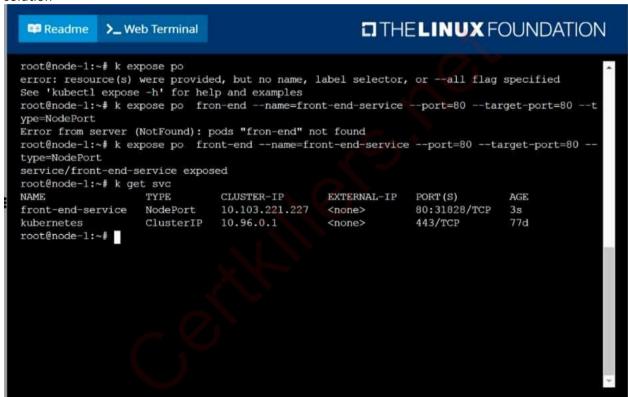
```
root@node-1:-# k create deploy nginx-app --image=nginx:1.11.10-alpine --dry-run=client -o y aml > app.yaml
root@node-1:-# vim app.yaml
root@node-1:-# k create -f app.yaml
deployment.apps/nginx-app created
root@node-1:-#
root@node-1:-# k set image deploy nginx-app nginx=nginx:1.11.13-alpine --record
deployment.apps/nginx-app image updated
root@node-1:-# k rollout undo deploy nginx-app
deployment.apps/nginx-app rolled back
root@node-1:-#
```

SIMULATION

Create and configure the service front-end-service so it's accessible through NodePort and routes to the existing pod named front-end.

Solution

solution



Question: 9

SIMULATION

Create a pod as follows:

- Name: mongo
- Using Image: mongo
- In a new Kubernetes namespace named: my-website

Solution

solution

```
THELINUX FOUNDATION
 Readme
            >_ Web Terminal
root@node-1:~#
root@node-1:~#
root@node-1:~# k create ns my-website
namespace/my-website created
root@node-1:~# k run mongo --image=mongo -n my-website
pod/mongo created
root@node-1:~# k get po -n my-website
       READY STATUS
0/1 ContainerCreating
                                  RESTARTS
                                            AGE
      0/1
                                            45
root@node-1:~#
```

Question: 10

SIMULATION

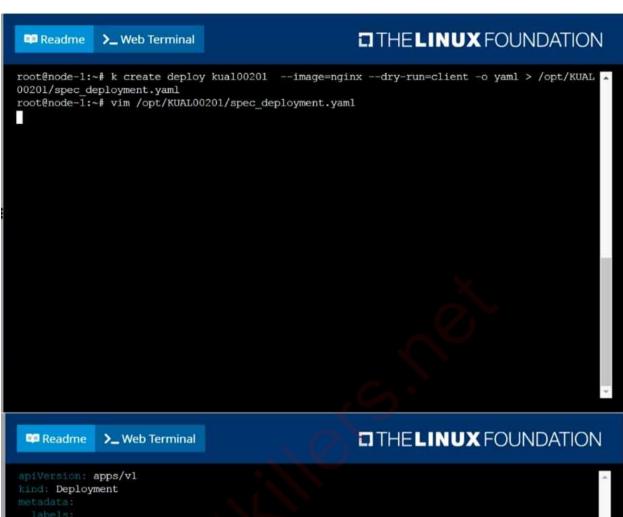
Create a deployment spec file that will:

- Launch 7 replicas of the nginx Image with the label app_runtime_stage=dev
- deployment name: kual00201

Save a copy of this spec file to /opt/KUAL00201/spec_deployment.yaml (or /opt/KUAL00201/spec_deployment.json).

When you are done, clean up (delete) any new Kubernetes API object that you produced during this task.

Solution



```
apiVersion: apps/vl
kind: Deployment
metadata:
labels:
    app_runtime_stage: dev
    name: kual00201
spec:
    replicas: 7
    selector:
    matchLabels:
    app_runtime_stage: dev
template:
    metadata:
    labels:
    app_runtime_stage: dev
spec:
    containers:
    - image: nginx
    name: nginx

"/opt/KUAL00201/spec_deployment.yaml" 19L, 320C written
```

SIMULATION

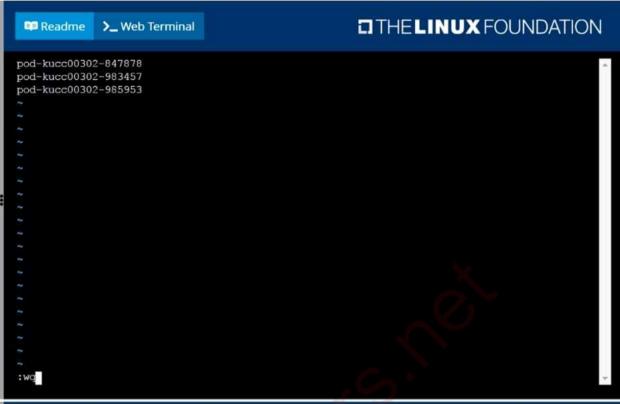
Create a file:

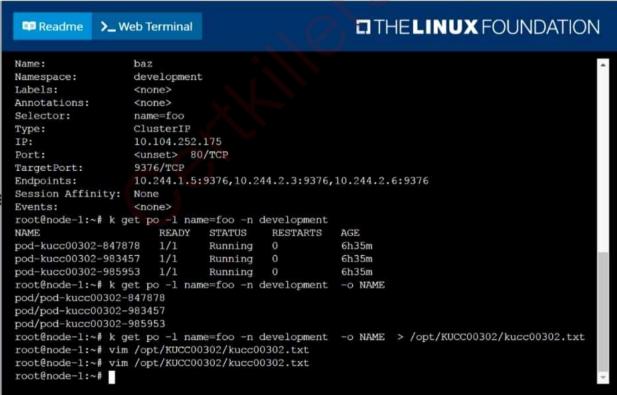
/opt/KUCC00302/kucc00302.txt that lists all pods that implement service baz in namespace development.

The format of the file should be one pod name per line.

Solution

```
THE LINUX FOUNDATION
 Readme
             >_ Web Terminal
root@node-1:~#
root@node-1:~# k describe svc baz -n development
                  baz
Namespace:
                  development
Labels:
                  <none>
Annotations:
                   <none>
                   name=foo
Selector:
                  ClusterIP
Type:
                  10.104.252.175
IP:
Port:
                  <unset> 80/TCP
                   9376/TCP
TargetPort:
Endpoints:
                  10.244.1.5:9376, 10.244.2.3:9376, 10.244.2.6:9376
Session Affinity: None
                   <none>
root@node-1:~# k get po -1 name=foo -n development
                      READY STATUS
pod-kucc00302-847878
                     1/1
                               Running
                                        0
                                                    6h35m
pod-kucc00302-983457 1/1
pod-kucc00302-985953 1/1
                               Running
                                                    6h35m
                                                    6h35m
                              Running
root@node-1:~# k get po -1 name=foo -n development -o NAME
pod/pod-kucc00302-847878
pod/pod-kucc00302-983457
pod/pod-kucc00302-985953
root@node-1:~# k get po -1 name=foo -n development -o NAME > /opt/KUCC00302/kucc00302.txt
root@node-1:~# vim /opt/KUCC00302/kucc00302.txt
```





SIMULATION

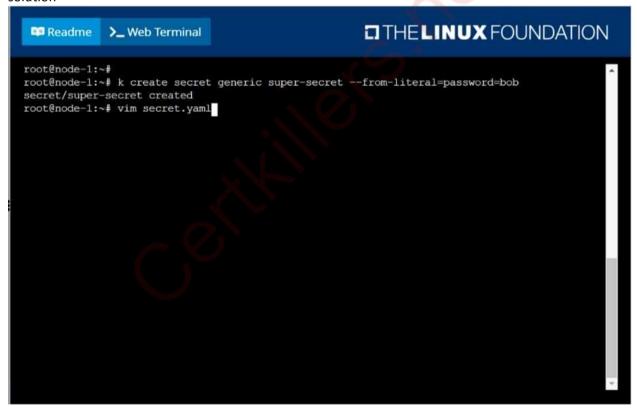
Create a Kubernetes secret as follows:

- Name: super-secret
- password: bob

Create a pod named pod-secrets-via-file, using the redis Image, which mounts a secret named super-secret at /secrets.

Create a second pod named pod-secrets-via-env, using the redis Image, which exports password as CONFIDENTIAL







THE LINUX FOUNDATION Readme >_ Web Terminal root@node-1:~# k create -f secret.yaml pod/pod-secrets-via-file created root@node-1:~# vim secret1.yaml root@node-1:~# k create -f secret1.yaml pod/pod-secrets-via-env created root@node-1:~# k get po NAME READY STATUS RESTARTS AGE cpu-utilizer-98b9se Running 6h25m cpu-utilizer-ab2d3s Running 6h25m 1/1 0 1/1 6h25m cpu-utilizer-kipb9a Running ds-kusc00201-2r2k9 1/1 Running 40m 0 ds-kusc00201-hzm9q 1/1 Running 40m Running 6h28m foo 1/1 6h27m front-end 1/1 Running Running hungry-bear 1/1 0 36m kucc8 3/3 Running 34m nginx-app-848cfcf495-9prjh 1/1 Running 19m nginx-app-848cfcf495-g12kh 1/1 Running 19m nginx-app-848cfcf495-pg2c8 1/1 Running 0 19m nginx-kusc00101 1/1 Running 26m pod-secrets-via-env 1/1 Running 0 4s pod-secrets-via-file 1/1 106s Running webserver-84c55967f4-qzjcv Running 1/1 0 6h43m webserver-84c55967f4-t4791 1/1 Running 6h43m root@node-1:~#

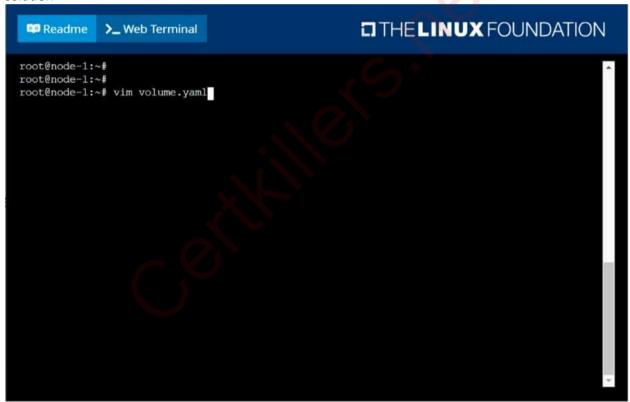
SIMULATION

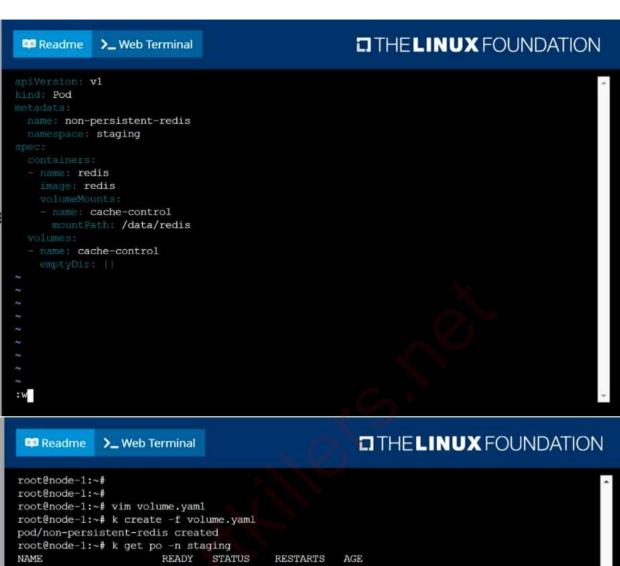
Create a pod as follows:

- Name: non-persistent-redis
- container Image: redis
- Volume with name: cache-control
- Mount path: /data/redis

The pod should launch in the staging namespace and the volume must not be persistent.

Solution



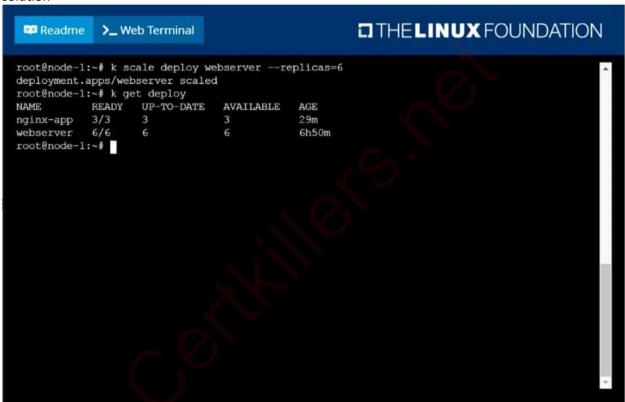


SIMULATION

Scale the deployment webserver to 6 pods.

Solution

solution



Question: 15

SIMULATION

Check to see how many worker nodes are ready (not including nodes tainted NoSchedule) and write the number to /opt/KUCC00104/kucc00104.txt.

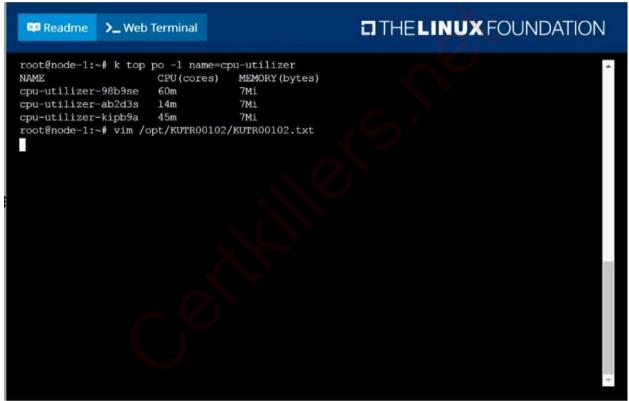
Solution



SIMULATION

From the pod label name=cpu-utilizer, find pods running high CPU workloads and write the name of the pod consuming most CPU to the file /opt/KUTR00102/KUTR00102.txt (which already exists).

Solution





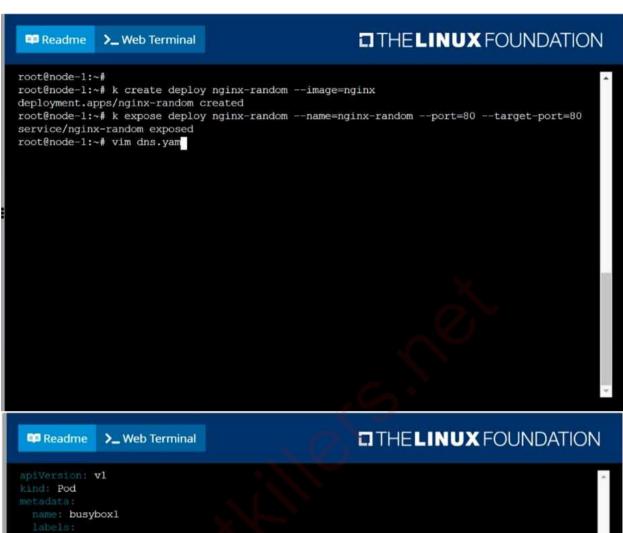
SIMULATION

Create a deployment as follows:

- Name: nginx-random
- Exposed via a service nginx-random
- Ensure that the service & pod are accessible via their respective DNS records
- The container(s) within any pod(s) running as a part of this deployment should use the nginx Image Next, use the utility nslookup to look up the DNS records of the service & pod and write the output to /opt/KUNW00601/service.dns and /opt/KUNW00601/pod.dns respectively.

Solution

Solution:



```
apiversion: v1
kind: Pod
metadata:
name: busybox
labels:
name: busybox
spec:
containers:
- sleep
- "3600"
name: busybox
```



SIMULATION

<u>Create a snapshot of the etcd instance running at https://127.0.0.1:2379, saving the snapshot to the file path /srv/data/etcd-snapshot.db.</u>

The following TLS certificates/key are supplied for connecting to the server with etcdctl:

- CA certificate: /opt/KUCM00302/ca.crt
- Client certificate: /opt/KUCM00302/etcd-client.crt
- Client key: Topt/KUCM00302/etcd-client.key

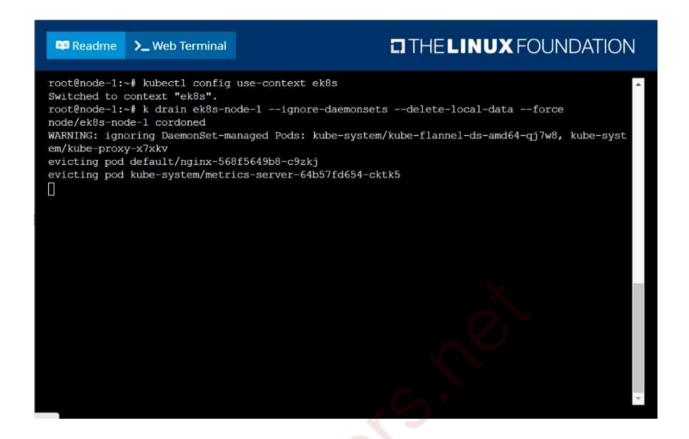
Solution



SIMULATION

Set the node named ek8s-node-1 as unavailable and reschedule all the pods running on it.

Solution



SIMULATION

A Kubernetes worker node, named wk8s-node-0 is in state NotReady. Investigate why this is the case, and perform any appropriate steps to bring the node to a Ready state, ensuring that any changes are made permanent.

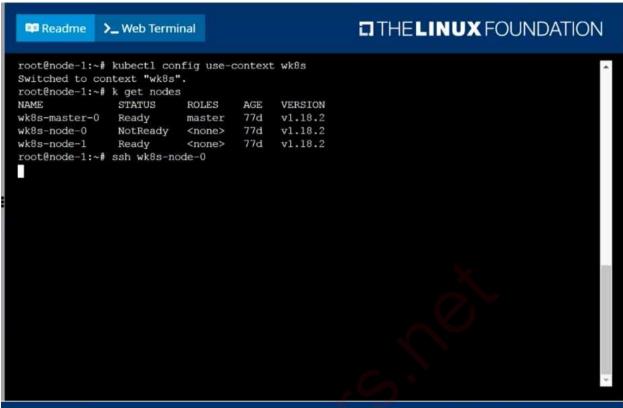
You can ssh to the failed node using:

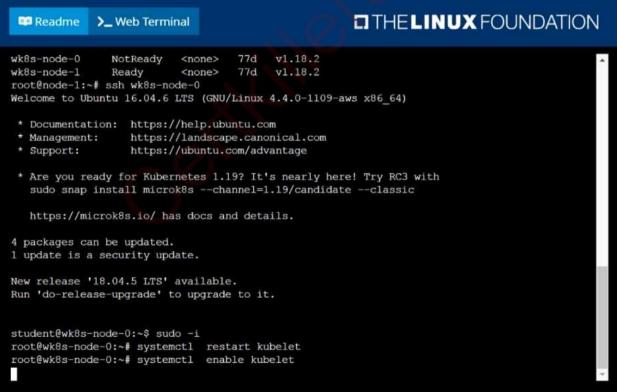
[student@node-1] \$ | ssh Wk8s-node-0

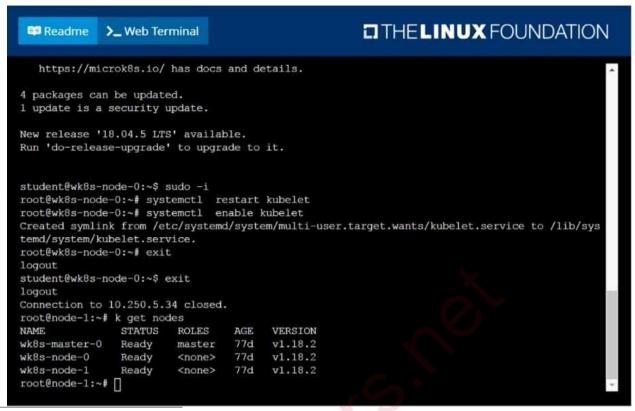
You can assume elevated privileges on the node with the following command:

[student@w8ks-node-0] \$ | sudo -i

Solution	







SIMULATION

Configure the kubelet systemd- managed service, on the node labelled with name=wk8s-node-1, to launch a pod containing a single container of Image httpd named webtool automatically. Any spec files required should be placed in the /etc/kubernetes/manifests directory on the node.

You can ssh to the appropriate node using:

[student@node-1] \$ ssh wk8s-node-1

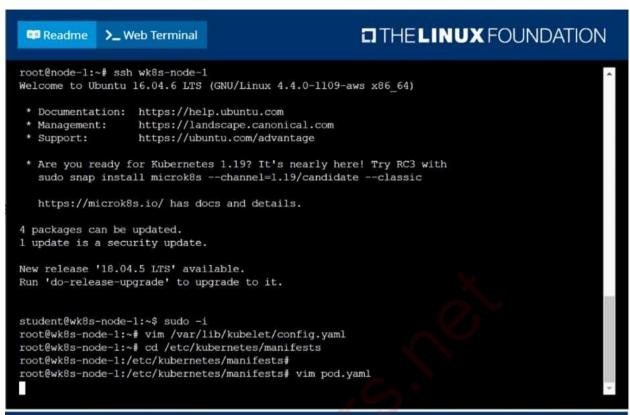
You can assume elevated privileges on the node with the following command:

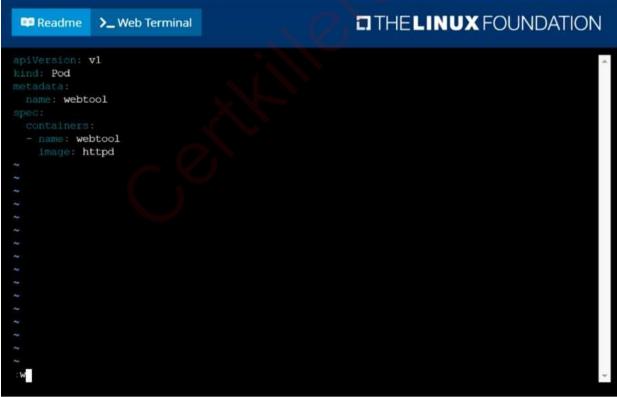
[student@wk8s-node-1] \$ | sudo -i

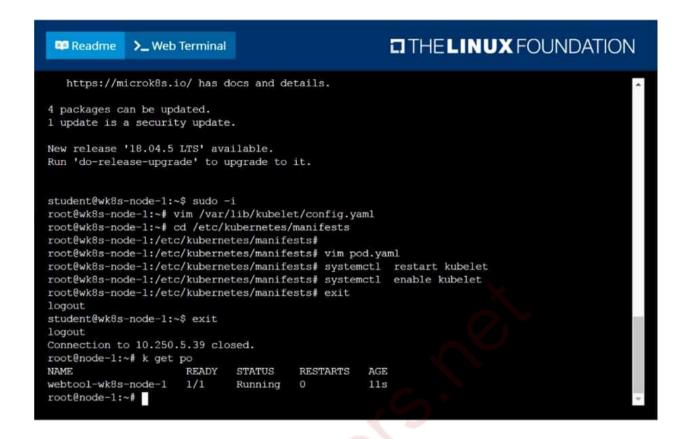
Solution

```
Readme
                                                    THE LINUX FOUNDATION
             >_ Web Terminal
root@node-1:~#
root@node-1:~# kubectl config use-context wk8s
Switched to context "wk8s".
root@node-1:~# ssh wk8s-node-1
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.4.0-1109-aws x86 64)
 * Documentation: https://help.ubuntu.com
 * Management: https://landscape.canonical.com
 * Support:
                  https://ubuntu.com/advantage
 * Are you ready for Kubernetes 1.19? It's nearly here! Try RC3 with
   sudo snap install microk8s --channel=1.19/candidate --classic
  https://microk8s.io/ has docs and details.
4 packages can be updated.
1 update is a security update.
New release '18.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.
student@wk8s-node-1:~$ sudo -i
root@wk8s-node-1:~# vim /var/lib/kubelet/config.yaml
```









SIMULATION

For this item, you will have to ssh to the nodes ik8s-master-0 and ik8s-node-0 and complete all tasks on these nodes. Ensure that you return to the base node (hostname: node-1) when you have completed this item.

Context

As an administrator of a small development team, you have been asked to set up a Kubernetes cluster to test the viability of a new application.

Task

You must use kubeadm to perform this task. Any kubeadm invocations will require the use of the -- ignore-preflight-errors=all option.

- Configure the node ik8s-master-O as a master node. .
- Join the node ik8s-node-o to the cluster.

Solution

solution

You must use the kubeadm configuration file located at /etc/kubeadm.conf when initializing your cluster.

You may use any CNI plugin to complete this task, but if you don't have your favourite CNI plugin's manifest URL at hand, Calico is one popular option:

https://docs.projectcalico.org/v3.14/manifests/calico.yaml

Docker is already installed on both nodes and apt has been configured so that you can install the required tools.

Question: 23

SIMULATION

Given a partially-functioning Kubernetes cluster, identify symptoms of failure on the cluster.

Determine the node, the failing service, and take actions to bring up the failed service and restore the health of the cluster. Ensure that any changes are made permanently.

You can ssh to the relevant I nodes (bk8s-master-0 or bk8s-node-0) using:

[student@node-1] \$ ssh <nodename>

You can assume elevated privileges on any node in the cluster with the following command:

[student@nodename] \$ | sudo –i

Solution

solution

```
THE LINUX FOUNDATION
 Readme
            >_ Web Terminal
root@node-1:~#
root@node-1:~# kubectl config use-context bk8s
Switched to context "bk8s".
root@node-1:~# ssh bk8s-master-0
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.4.0-1109-aws x86 64)
* Documentation: https://help.ubuntu.com
 * Management:
                  https://landscape.canonical.com
                  https://ubuntu.com/advantage
 * Support:
 * Are you ready for Kubernetes 1.19? It's nearly here! Try RC3 with
  sudo snap install microk8s --channel=1.19/candidate --classic
  https://microk8s.io/ has docs and details.
4 packages can be updated.
1 update is a security update.
New release '18.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.
student@bk8s-master-0:~$ sudo -i
root@bk8s-master-0:~# vim /var/lib/kubelet/config.yaml
```



Question: 24

root@node-1:~#

root@bk8s-master-0:~#
root@bk8s-master-0:~# exit

student@bk8s-master-0:~\$ exit

Connection to 10.250.4.77 closed.

SIMULATION

Create a persistent volume with name app-data, of capacity 2Gi and access mode ReadWriteMany. The type of volume is hostPath and its location is /srv/app-data.

Solution

solution

Persistent Volume

A persistent volume is a piece of storage in a Kubernetes cluster. PersistentVolumes are a cluster-level resource like nodes, which don't belong to any namespace. It is provisioned by the administrator and has a particular file size. This way, a developer deploying their app on Kubernetes need not know the underlying infrastructure. When the developer needs a certain amount of persistent storage for their application, the system administrator configures the cluster so that they consume the PersistentVolume provisioned in an easy way.

Creating Persistent Volume kind: Persistent Volume apiVersion: v1 metadata:

name:app-data

spec:

capacity: # defines the capacity of PV we are creating storage: 2Gi #the amount of storage we are tying to claim accessModes: # defines the rights of the volume we are creating

- ReadWriteMany

hostPath:

path: "/srv/app-data" # path to which we are creating the volume

Challenge

• Create a Persistent Volume named app-data, with access mode ReadWriteMany, storage classname shared, 2Gi of storage capacity and the host path /srv/app-data.

```
apiVersion: v1
kind: PersistentVolume
 name app-data
    storage 2Gi

    ReadWriteManv

      path:/srv/app-data
 storageClassName: shared
"app-data.yaml" 12L, 194C
```

2. Save the file and create the persistent volume.

njerry191@cloudshell:~ (extreme-clone-265411)\$ kubectl create -f pv.yaml persistentvolume/pv created 3. View the persistent volume.

```
njerry191@cloudshell:~ (extre
                                  clone-265411) $ kubectl get pv
                                  RECLAIM POLICY
                                                                         STORAGECLASS
                                                    STATUS
                                                    Available
```

• Our persistent volume status is available meaning it is available and it has not been mounted yet. This status will change when we mount the persistentVolume to a persistentVolumeClaim.

PersistentVolumeClaim

In a real ecosystem, a system admin will create the PersistentVolume then a developer will create a PersistentVolumeClaim which will be referenced in a pod. A PersistentVolumeClaim is created by specifying the minimum size and the access mode they require from the persistentVolume.

• Create a Persistent Volume Claim that requests the Persistent Volume we had created above. The claim should request 2Gi. Ensure that the Persistent Volume Claim has the same storageClassName as the persistentVolume you had previously created.

kind: PersistentVolume

apiVersion: v1

metadata:

name:app-data

spec:

accessModes:

- ReadWriteMany

resources:

requests:

storage: 2Gi

storageClassName: shared

2. Save and create the pvc

njerry191@cloudshell:~ (extreme-clone-2654111)\$ kubect1 create -f app-data.yaml persistentvolumeclaim/app-data created

3. View the pvc

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS
pv Bound pv 512m RWX shared
```

4. Let's see what has changed in the pv we had initially created.

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS REASON AGE
pv 512m RWX Retain Bound default/pv shared 16m
```

Our status has now changed from available to bound.

5. Create a new pod named myapp with image nginx that will be used to Mount the Persistent Volume Claim with the path /var/app/config.

Mounting a Claim apiVersion: v1 kind: Pod metadata:

creationTimestamp: null

name: app-data

spec:

volumes:

name:congigpvc persistenVolumeClaim: claimName: app-data

containers:
- image: nginx
name: app
volumeMounts:

- mountPath: "/srv/app-data"

name: configpvc

Question: 25

Create a namespace called 'development' and a pod with image nginx called nginx on this namespace.

kubectl create namespace development kubectl run nginx --image=nginx --restart=Never -n development

Question: 26

Create a nginx pod with label env=test in engineering namespace

	Solution
kubectl run nginximage=nginxrestart=Neverlabels=env=testna	amespace=engineeringdry-
run -o yaml > nginx-pod.yaml kubectl run nginximage=nginxrestart=Neverlabels=env=testna run -o yaml kubectl create -n engineering -f — YAML File:	amespace=engineeringdry-
apiVersion: v1	
kind: Pod	
metadata:	
name: nginx	
namespace: engineering	
labels:	
env: test	
spec:	
containers:	
- name: nginx	
image: nginx	
imagePullPolicy: IfNotPresent	
restartPolicy: Never	
kubectl create -f nginx-pod.yaml	
Question: 27	
<u> </u>	
Get list of all pods in all namespaces and write it to file "/opt/pods-list.ya	aml"
det list of all pous in all numespaces and write it to life yopi, pous list.y.	31111
	Solution
	Solution
kubectl get po –all-namespaces > /opt/pods-list.yaml	
Question: 28	
Create a pod with image nginx called nginx and allow traffic on port 80	
	Solution
kubectl run nginximage=nginxrestart=Neverport=80	
made in inches i	
Question: 29	
Question: 29	

Create a busybox pod that runs the command "env" and save the output to "envpod" file

	Solution
kubectl run busyboximage=busyboxrestart=Never —-rm -it env > e	nvpod.yaml
Question: 30	
List pod logs named "frontend" and search for the pattern "started" and logs"	d write it to a file "/opt/error-
	Solution
Kubectl logs frontend grep -i "started" > /opt/error-logs	
Question: 31	
Create a pod that echo "hello world" and then exists. Have the pod de completed	leted automatically when it's
	Solution
kubectl run busyboximage=busybox -itrmrestart=Never/bin/sh -c 'echo hello world'	
kubectl get po # You shouldn't see pod with the name "busybox"	
Question: 32	
Create a pod with environment variables as var1=value1.Check the environment	ronment variable in pod
	Solution
kubectl run nginximage=nginxrestart=Neverenv=var1=value1 # then	
kubectl exec -it nginx env # or	
kubectl exec -it nginx sh -c 'echo \$var1' # or	
kubectl describe po nginx grep value1	

Question: 33	
Get list of all the pods showing name and namespace with a jsonpath ex	xpression.
	Solution
kubectl get pods -o=jsonpath="{.items[*]['metadata.name' , 'metadata.namespace']}"	
Question: 34	
Check the image version in pod without the describe command	O'
	Solution
kubectl get po nginx -o jsonpath='{.spec.containers[].image}{"\n"}'	
Question: 35	
List the nginx pod with custom columns POD_NAME and POD_STATUS	
	Solution
kubectl get po -o=custom-columns="POD_NAME:.metadata.name, POD_STATUS:.status.containerStatuses[].state"	
Question: 36	
List all the pods sorted by name	
	Solution
kubect1 get podssort-by=.metadata.name	
Question: 37	

Create a r	ood that	having 3	containers in i	it? ((Multi-Container)

Solution

image=nginx, image=redis, image=consul

Name nginx container as "nginx-container"

Name redis container as "redis-container"

Name consul container as "consul-container"

Create a pod manifest file for a container and append container

section for rest of the images

kubectl run multi-container --generator=run-pod/v1 --image=nginx --

dry-run -o yaml > multi-container.yaml

then

vim multi-container.yaml

apiVersion: v1 kind: Pod metadata: labels:

run: multi-container name: multi-container

spec: containers: - image: nginx

name: nginx-container

- image: redis

name: redis-container - image: consul

name: consul-container restartPolicy: Always

Question: 38

Create 2 nginx image pods in which one of them is labelled with env=prod and another one labelled with env=dev and verify the same.

Solution

kubectl run --generator=run-pod/v1 --image=nginx -- labels=env=prod nginx-prod --dry-run -o yaml > nginx-prodpod.yaml Now, edit nginx-prod-pod.yaml file and remove entries like "creationTimestamp: null" "dnsPolicy: ClusterFirst"

vim nginx-prod-pod.yaml

apiVersion: v1 kind: Pod metadata: labels: env: prod name: nginx-prod spec: containers: - image: nginx name: nginx-prod restartPolicy: Always # kubectl create -f nginx-prod-pod.yaml kubectl run --generator=run-pod/v1 --image=nginx -labels=env=dev nginx-dev --dry-run -o yaml > nginx-dev-pod.yaml apiVersion: v1 kind: Pod metadata: labels: env: dev name: nginx-dev spec: containers: - image: nginx name: nginx-dev restartPolicy: Always # kubectl create -f nginx-prod-dev.yaml Verify: kubectl get po --show-labels kubectl get po -l env=prod kubectl get po -l env=dev **Question: 39** Get IP address of the pod - "nginx-dev" **Solution** Kubect1 get po -o wide Using JsonPath kubect1 get pods -o=jsonpath='{range $. items[*]] {.metadata.name} { "\t"} {.status.podIP} { "\n"} { end} { "} { end} { "} { end} { end}$ Question: 40 Print pod name and start time to "/opt/pod-status" file Solution kubect1 get pods -o=jsonpath='{range

.items[*]]{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'

Question: 41		
Check the Image version of nginx-dev p	ood using jsonpath	Solution
<pre>kubect1 get po nginx-dev -o jsonpath='{.spec.containers[].image}{"'</pre>	\n"}'	
Question: 42		
Create a busybox pod and add "sleep 3	600" command	Solution
kubectl run busyboximage=busybox - "sleep 3600"	-restart=Never /bin/sh -c	
Question: 43		
Create an nginx pod and list the pod wi	th different levels of verbosity	Calintian
// create a pod kubectl run nginximage=nginxresta // List the pod with different verbosity kubectl get po nginxv=7 kubectl get po nginxv=8 kubectl get po nginxv=9 Question: 44	rt=Neverport=80	Solution
List the nginx pod with custom column	s POD_NAME and POD_STATUS	Solution
kubectl get po -o=custom-columns="POD_STATUS:.status.containerStatuses	_	

kubectl delete po "nginx-prod"

Question: 45	
List all the pods sorted by name	
-	
-	Solution
kubectl get podssort-by=.metadata.name	
Question: 46	
List all the pods sorted by created timestamp	o'\
	Solution
kubect1 get podssort-by=.metadata.creationTimestamp	
Question: 47	
List all the pods showing name and namespace with a json path expressi	on
	Solution
kubectl get pods -o=jsonpath="{.items[*]['metadata.name', 'metadata.namespace']}"	
Question: 48	
List "nginx-dev" and "nginx-prod" pod and delete those pods	
	Solution
luibact at and a cuida	
kubect1 get pods -o wide kubectl delete po "nginx-dev"	